GLIMMER-MG Release Notes
Version 0.1

David R. Kelley, Art L. Delcher

17 May 2011

1 Introduction

This document describes Version 0.1 of the GLIMMER-MG gene-finding software for metagenomic sequences. Users discovering problems or errors are encouraged to report them to `dakelley@umiacs.umd.edu`.

GLIMMER is a collection of programs for identifying genes in microbial DNA sequences. The system works by creating a variable-length Markov model from a training set of genes and then using that model to attempt to identify all genes in a given DNA sequence. The three versions of GLIMMER are described in [SDKW98], [DHK⁺99], and [DBPS07].

GLIMMER-MG is released as OSI Certified Open Source Software under the Artistic License. The license is contained in the file, `LICENSE`, in the distribution.

2 Installation

GLIMMER-MG software was written for the Linux software environment. The following instructions assume a Linux system. They also work under Mac OSX.

To install GLIMMER-MG, download the compressed tarfile `glimmer-mg-1.0.1.tar.gz` from the website. Then uncompress the file by typing

```
tar xzf glimmer-mg.1.0.1.tar.gz
```

A directory named `glimmer-mg` should result. In that directory, is a subdirectory named `src`. Within the `src` subdirectory type

```
make
```

(or alternately `gmake`). This will compile the GLIMMER-MG programs and put the executable files in the directory `glimmer-mg/bin`. These files can be copied or moved to whatever directory is convenient to the user.

3 Running Glimmer-MG

GLIMMER-MG can be run in a few different modes, depending on the characteristics of the sequences provided, e.g., whether they represent accurate contigs or reads from the Illumina or 454 technologies. Then within each mode, GLIMMER-MG can be run with varying amounts of classification/clustering preprocessing.

3.1 Preprocessing options

The GLIMMER-MG pipeline can be run a number of different ways, including a Python script to run the full pipeline and options to classify/cluster the reads separately and predict genes using those results.

3.1.1 Running from scratch

The Python script `glimmer-mg.py` runs the full Glimmer-MG pipeline.

```
glimmer-mg.py seqs.fasta
```

This script will classify the sequences with Phymm [BS09], make initial predictions for all sequences, cluster the sequences with Scimm [KS10], and make final predictions within each cluster.

3.1.2 Running from scratch with no retraining

By passing the option `--iter 0` to `glimmer-mg.py`, the clustering and re-training steps will be skipped. That is, the sequences will be classified with Phymm and predictions made.

```
glimmer-mg.py --iter 0 seqs.fasta
```

3.1.3 Classification separate

Some users may prefer to run the computationally intensive classification of sequences separately (e.g. on a computer cluster). `glimmer-mg.py` can be made to expect this by using the `--raw` option, which specifies that the raw Phymm output file already exists in the current working directory.

```
glimmer-mg.py --raw seqs.fasta
```

3.1.4 Clustering separate

In a similar vein, clustering can be performed separately before making gene predictions by specifying the `--clust` option, which specifies that Scimm or PhyScimm output is in the current working directory.

```
glimmer-mg.py --clust seqs.fasta
```

3.2 Sequence modes

Depending on the error characteristics of the input sequences, Glimmer-MG has three modes that best handle certain types of data. In each case, the options described can be passed to the C++ binary `glimmer-mg` or the Python script to manage the entire pipeline `glimmer-mg.py`.

3.2.1 Accurate contigs

If you believe that the sequences on which you would like to find genes are very accurate, than run Glimmer-MG in the default mode with no additional options.

3.2.2 454 indels

If the input sequences are reads or contigs from the 454 or Ion Torrent technologies and indel errors are expected to exist, GLIMMER-MG can be made to predict such errors and account for them in its gene predictions by using *indel* mode.

```
glimmer-mg.py --indels 454.fasta
```

3.2.3 Illumina substitution errors

If the input sequences are reads or contigs from the Illumina technology and substitution errors are expected to exist, GLIMMER-MG can be made to predict such errors and account for them in its gene predictions by using *substitution* mode.

```
glimmer-mg.py --sub illumina.fasta
```

4 Sample Run Directory

A directory containing a sample run of GLIMMER3 is provided. This directory, named **sample_run** contains the genome sequence for *Treponema pallidum* (file **tpall.fna**) and a list of annotated genes for it (file **tpall.nh**), both downloaded from GenBank. The files whose names begin **from-scratch** are the result of running the script

```
g3-from-scratch.csh tpall.fna from-scratch
```

The files whose names begin **from-training** are the result of running the script

```
g3-from-training.csh tpall.fna tpall.nh from-training
```

The files whose names begin **iterated** are the result of running the script

```
g3-iterated.csh tpall.fna iterated
```

Users will need to modify the path directories at the top of these scripts to be able to run them (see Section **??** above).

5 Notes on Glimmer-MG programs

5.1 `glimmer-mg` Program

This is the main program that makes gene preditions.

5.1.1 `glimmer-mg` Parameters & Options

The invocation for `glimmer-mg` is:

```
glimmer-mg [options] sequence tag
```

where *sequence* is the name of the file containing the DNA sequence(s) to be analyzed and *tag* is a prefix used to name the output files.

options can be the following:

-b *filename* or **--rbs_pwm** *filename*

Read a position weight matrix (PWM) from *filename* to identify the ribosome binding site to help choose start sites. The format of this file is indicated by the following example:

```
6
a      212       309        43        36       452       138
c       55        58         0        19        48        26
g      247       141       501       523         5       365
t       64        70        34         0        73        49
```

The first line is the number of positions in the pattern, *i.e.*, the number of columns in the matrix (not counting the first column of labels). The column values are the relative frequencies of nucleotides at each position.

-c *filename* or **--class** *filename*

Read sequence classifications from *filename*. Each line of *filename* should specify a sequence fasta header followed by a series of GenBank reference genome classifications. The classifications should name the reference genome's Phymm directory (corresponding to the GenBank strain name), followed by the character '|', followed by the genomic sequence's unique identifier. It is our expectation that users will not create these files, but they generate them with **glimmer-mg.py**. For example:

```
read1      Mycobacterium_leprae_Br4923|NC_011896 Corynebacterium_glutamicum_ATCC_13032|NC_003450
read2      Rhodopseudomonas_palustris_BisA53|NC_008435 Rhodopseudomonas_palustris_HaA2|NC_007778
...
```

-f *filename* or **--features** *filename*

filename specifies counts for features such as gene length, start codon usage, adjacent gene orientations, and adjacent gene distances for coding and noncoding ORFs. GLIMMER-MG will convert these counts into probability models for use in loglikelihood computations for the features. It is our expecation that users will not create these files, but generate them with **train_features.py** (via **glimmer-mg.py**). For an example of the input format, see the **sample_run** directory.

-g *n* or **--gene_len** *n*

Set the minimum gene length to *n* nucleotides. This does not include the bases in the stop codon.

-h or **--help**

Print the usage message.

-i or **--indel**

Predict genes in "indel-mode" where gene predictions may shift the coding frame, implicitly predicting an insertion or deletion in the sequence. If quality values are provided with the **-q** option, positions at which frame shifts should be considered will be identifed by their low quality. If no quality values are provided, frame shifts will be considered at long homopolymer runs, assuming that the sequences were generated by the 454 or Ion torrent technologies.

-m *filename* or **--icm** *filename*

filename contains an ICM trained on coding sequence to be used for ORF scoring rather than ICM's obtained via classification results.

-o *n* or **--max_olap** *n*

Set the maximum overlap length to *n*. Overlaps of this many or fewer bases are allowed between genes. The new dynamic programming algorithm should <u>never</u> output genes that overlap by more than this many bases.

-q *filename* or **--quality** *filename*

filename contains quality values in fasta format matching up with the sequences fasta file. The quality values are used in "indel-mode" and "substitution-mode" to identify low quality positions in the sequences where errors are most likely.

-r or **--circular**

Assume a circular rather than linear genome, *i.e.*, there may be genes that "wraparound" between the beginning and end of the sequence.

-s or **--sub**

Predict genes in "substitution-mode" where gene predictions may pass through stop codons, implicitly predicting a sequencing error that mutated a standard codon to a stop codon. If quality values are provided with the **-q** option, they will be used to compute a log-likelihood penalty for passing through a given stop codon. Otherwise, we assume all quality values are Phred 30.

-u *n* or **--fudge** *n*

n specifices a "fudge factor" to be added to the log-likeihood score for every ORF. The "fudge factor" acts as a means to tune the sensitivity versus specificity of the predictions.

-z *n* or **--trans_table** *n*

Use Genbank translation table number *n* to specify stop codons.

-Z *codon-list* or **--stop_codons** *codon-list*

Specify stop codons as a comma-separated list. Sample format: **-Z tag,tga,taa**. The default stop codons are **tag**, **tga** and **taa**.

5.1.2 Glimmer-MG Output Formats

`.predict` **File**

This file has the final gene predictions. It's format is the fasta-header line of the sequence followed by one line per gene. Here is a sample of the beginning of such a file:

```
>SRR029690.117009 E4LJNJL02B7ZT9 length=531
orf00045       386         -1  -2    52.14 I:272 D: S:
orf00057       532        450  -3     3.86 I: D:507 S:
```

The columns are:

Column 1   The identifier of the predicted gene. The numeric portion matches the number in the `ID` column of the `.detail` file.

Column 2   The start position of the gene.

Column 3   The end position of the gene. This is the last base of the stop codon, *i.e.*, it includes the stop codon.

Column 4   The reading frame.

Column 5   The log-likelihood ratio score of the gene.

6 Notes on Glimmer programs

6.1 `build-icm` **Program**

This program constructs an interpolated context model (ICM) from an input set of sequences.

6.1.1 `build-icm` **Parameters & Options**

The format for invoking `build-icm` is:

> `build-icm` [*options*] *output-file* `<` *input-file*

Sequences are reads from standard input, the ICM is built and written to *output-file*. If *output-file* is "-", then the output will be sent to standard output. Since input comes from standard input, one also can "pipe" the input into this program, *e.g.*,

```
cat abc.in | build-icm xyz.icm
```

or even type in the input directly.

Possible *options* are:

-d *num*   or   **--depth** *num*

> Set the depth of the ICM to *num*. The depth is the maximum number of positions in the context window that will be used to determine the probability of the predicted position. The default value is 7.

-F   or   **--no_stops**

> Do not use any input strings with in-frame stop codons. Stop codons are determined by either the **-z** or **-Z** option.

-h   or   **--help**

> Print the usage message.

-p *num*   or   **--period** *num*

> Set the period of the ICM to *num*. The period is the number of different submodels for different positions in the text in a cyclic pattern. *E.g.*, if the period is 3, the first submodel will determine positions $1, 4, 7, \ldots$; the second submodel will determine positions $2, 5, 8, \ldots$; and the third submodel will determine positions $3, 6, 9, \ldots$. For a non-periodic model, use a value of 1. The default value is 3.

-r   or   **--reverse**

> Use the reverse of the input strings to build the ICM. Note that this is merely the reverse and <u>*NOT*</u> the reverse-complement. In other words, the model is built in the backwards direction.

-t   or   **--text**

> Output the model in a text format. This is for informational/debugging purposes only—the `glimmer3` program cannot read models in this form.

> The format of the output is a header line containing the parameters of the model, followed by individual probability lines. The entries on each probability line are:

> | Column | Description |
> |:------:|:------------|
> | 1 | ID number |
> | 2 | Context pattern |
> | 3 | Mutual information |
> | 4 | Probability of A |
> | 5 | Probability of C |
> | 6 | Probability of G |
> | 7 | Probability of T |

> The context pattern is divided into codons by the vertical lines (this option assumes the default 3-periodic model). The "?" represents the position being predicted. Letters represent specific values in their respective positions in the context window. The asterisk indicates the position that has maximum mutual information with the predicted position.

-v *num* or **--verbose** *num*

>   Set the verbose level to *num*. This controls extra debugging output—the higher the
>   value the more output.

-w *num* or **--width** *num*

>   Set the width of the ICM to *num*. The width includes the predicted position. The
>   default value is 12.

-z *n* or **--trans_table** *n*

>   Use Genbank translation table number *n* to specify stop codons.

-Z *codon-list* or **--stop_codons** *codon-list*

>   Specify stop codons as a comma-separated list. Sample format: `-Z tag,tga,taa`.
>   The default stop codons are `tag`, `tga` and `taa`.

6.2 `long-orfs` Program

This program identifies long, non-overlapping open reading frames (orfs) in a DNA se-
quence file. These orfs are very likely to contain genes, and can be used as a set of training
sequences for the `build-icm` program. More specifically, among all orfs longer than a min-
imum length ℓ, those that do not overlap any others are output. The start codon used for
each orf is the first possible one. The program, by default, automatically determines the
value ℓ that maximizes the number of orfs that are output. With the **-t** option, the initial
set of candidate orfs also can be filtered using entropy distance, which generally produces
a larger, more accurate training set, particularly for high-GC-content genomes. Entropy
distance is described in [OZWS04].

6.2.1 `long-orfs` Parameters & Options

The format for invoking `long-orfs` is:

>   `long-orfs` [*options*] *sequence* *output*

where *sequence* is the name of the file containing the DNA sequence to be analyzed and
output is the name of the output file of coordinates. *sequence* may contain only one se-
quence. If *output* is "**-**", then the output is directed to standard output.

>   Possible *options* are:

-A *codon-list* or **--start_codons** *codon-list*

>   Specify allowable start codons as a comma-separated list. Sample format:
>   `-A atg,gtg`. The default start codons are `atg`, `gtg` and `ttg`.

-E *filename* or **--entropy** *filename*

>   Read entropy profiles from *filename*. The format is one header line, then 20 lines of
>   3 columns each, which is the format produced by the program `entropy-profile` with

the **-b** option. The columns are amino acid, positive entropy, and negative entropy, respectively. Rows must be in alphabetical order by amino acid code letter.

The entropy profiles are used only if the **-t** option is specified.

-f or **--fixed**

Do *NOT* automatically calculate the minimum gene length that maximizes the number or length of output regions, but instead use either the value specified by the **-g** option or else the default, which is 90.

-g n or **--min_len** n

Set the minimum gene length to n nucleotides. This does not include the bases in the stop codon.

-h or **--help**

Print the usage message.

-i *filename* or **--ignore** *filename*

filename specifies regions of bases that are off limits, so that no bases within that area will be examined. The format for entries in this file is described above for the same option in the **glimmer3** program.

-l or **--linear**

Assume a linear rather than circular genome, *i.e.*, there will be no "wraparound" genes with part at the beginning of the sequence and the rest at the end of the sequence.

-L or **--length_opt**

Find and use as the minimum gene length the value that maximizes the total *length* of non-overlapping genes, instead of the default behaviour, which is to maximize the total *number* of non-overlapping genes.

-n or **--no_header**

Do not include the program-settings header information in the output file. With this option, the output file will contain only the coordinates of the selected orfs.

-o n or **--max_olap** n

Set the maximum overlap length to n. Overlaps of this many or fewer bases between genes are not regarded as overlaps.

-t x or **--cutoff** x

Only genes with an entropy distance score less than x will be considered. This cutoff is made before any subsequent steps in the algorithm.

-w or **--without_stops**

Do *NOT* include the stop codon in the region described by the output coordinates. By default it is included.

`-z` *n*   or   `--trans_table` *n*

    Use Genbank translation table number *n* to specify stop codons.

`-Z` *codon-list*   or   `--stop_codons` *codon-list*

    Specify allowable stop codons as a comma-separated list. Sample format:
    `-Z tag,tga`. The default stop codons are `tag`, `tga` and `taa`.

7   Versions

7.1   Version 0.10

- Initial release.

References

[BS09]   A. Brady and S.L. Salzberg. Phymm and phymmbl: metagenomic phylogenetic classification with interpolated markov models. *Nature Methods*, 6(9):673–676, 2009.

[DBPS07]   A.L. Delcher, K.A. Bratke, E.C. Powers, and S.L. Salzberg. Identifying bacterial genes and endosymbiont dna with glimmer. *Bioinformatics*, 23(6):673, 2007.

[DHK$^+$99]   A. Delcher, D. Harmon, S. Kasif, O. White, and S. Salzberg. Improved microbial gene identification with GLIMMER. *Nucl. Acids Res.*, 27(23):4636–4641, 1999.

[KS10]   D.R. Kelley and S.L. Salzberg. Clustering metagenomic sequences with interpolated markov models. *BMC Bioinformatics*, 11(1):544, 2010.

[OZWS04]   Z. Ouyang, H. Zhu, J. Wang, and S.Z. She. Multivariate entropy distance method for prokaryotic gene identification. *J. Bioinformatics & Comp. Biol.*, 2(2):353–373, 2004.

[SDKW98]   S. Salzberg, A. Delcher, S. Kasif, and O. White. Microbial gene identification using interpolated Markov models. *Nucl. Acids Res.*, 26(2):544–548, 1998.